

Reflection on Global AI Governance and Suitable Paths Forward

The surge of generative AI in late 2022 created a moment unlike anything I have seen before in computing. Overnight, it seemed, tools capable of producing text, images, code, and even complex reasoning became accessible to anyone with an internet connection. Although AI research has existed for decades, this wave of generative systems feels qualitatively different. The scale, speed, and unpredictability of their development have forced societies, industries, and governments into conversations that they were not fully prepared for. As someone studying and working in the field, I find myself oscillating between excitement about the possibilities and concern about the lack of cohesive governance. Reading Correa et al. (2023) helped me frame these mixed feelings by showing just how fragmented global AI ethics and regulation really are.

Correa et al. argue that although many organisations and governments are trying to articulate the values that should guide AI development—such as transparency, fairness, accountability, and human oversight—there is little consensus about what these values mean in practice. Their study compares governance documents from around the world and highlights how cultural, political, and economic differences shape interpretations of ethical principles. This point struck me because it explains why attempts to “standardise” AI ethics often feel vague or overly aspirational. Even when countries use the same ethical vocabulary, their priorities diverge once implementation is considered. A term like “responsibility,” for example, may imply strict regulatory oversight in one jurisdiction but voluntary best-practice guidelines in another. Reflecting on this made me realise that one of the biggest governance challenges is not the lack of ethical principles but rather the difficulty of operationalising them across such diverse contexts.

Deckard (2023) approaches the issue from a different angle, focusing on the pace at which the private sector is adopting and deploying generative AI. The article captures the intense commercial pressure behind rapid innovation and the way companies are compelled to release increasingly capable models in order to compete. This acceleration leaves regulators struggling to keep up, and even computing professionals often feel disoriented by how quickly expectations shift. When I reflect on my own experience, I recognise this sense of constant movement: new tools arrive before we fully understand the consequences of the previous ones; workflows evolve before policies catch up; and social norms lag behind the technological realities. Deckard’s argument underscores the urgency of building governance systems that are adaptable rather than static.

Combining insights from both papers shaped my view that the most realistic and ethically responsible path forward is a *multi-layered governance model*. On a global scale, I believe there must be some shared, foundational principles—simple enough to gain wide agreement but strong enough to protect against the most serious harms. These principles might include commitments to safety, traceability, and accountability. Safety, in this context, should not be a one-off checkpoint but an ongoing process of evaluation and monitoring, especially for models that evolve over time or can be fine-tuned for unpredictable purposes. Traceability, or at least some level of transparency, is essential so that users and regulators can understand how an AI system works, what data it was trained on, and where its limitations lie.

Accountability ensures that responsibility does not evaporate into the complexity of the technology, leaving individuals or communities to bear the consequences alone.

However, while I support global coordination, I also recognise that governance cannot be identical everywhere. Countries have different political systems, economic incentives, and cultural values, all of which influence how they perceive both risks and opportunities in AI. The European Union, for instance, tends to prioritise human rights and regulatory oversight; the United States leans more heavily on innovation and market-driven solutions; China integrates AI governance into broader goals of social stability and state planning. Expecting these models to converge into a single unified framework is unrealistic. Instead, I think a more effective approach is one where international bodies establish shared minimum standards—perhaps through an institution similar to those that govern aviation or data privacy—while allowing national governments to build detailed regulations that reflect local priorities.

Reflecting on these ideas also made me consider their implications for legal, social, and professional issues. Legally, harmonised international standards could reduce uncertainty for companies operating across borders and provide clearer avenues for citizens seeking redress when harm occurs. Socially, stronger governance could help rebuild public trust, which has been weakened by concerns about misinformation, bias, job displacement, and the opacity of AI systems. Without trust, even beneficial AI innovations risk being rejected or misused. Professionally, the responsibilities placed on computing practitioners would become more explicit. Codes of ethics from organisations like ACM or BCS already emphasise duties such as ensuring transparency, protecting privacy, and prioritising societal welfare, but these principles gain more weight when embedded into regulatory structures. Reflecting on this reminded me that technical expertise alone is not enough; computing professionals must understand the societal implications of their design choices and be prepared to justify them.

Ultimately, this reflection has led me to a position that balances optimism with caution. I believe generative AI offers immense potential to support education, creativity, science, and industry. Yet I also see how quickly poorly governed technologies can undermine public institutions, distort information ecosystems, or exacerbate inequalities. A multi-layered approach to governance—international coordination paired with national flexibility—seems to me both practical and ethically grounded. It acknowledges the global nature of AI while respecting the realities of political diversity. It also places responsibility not only on governments but on industry and individual professionals who design, deploy, and interact with these systems.

Writing this essay has helped me better understand the interconnected roles of regulation, culture, professional ethics, and technological innovation. AI is no longer a niche research field; it has become a central force shaping the future of work, communication, and societal structures. The way we choose to govern it now will determine not just the safety and reliability of future systems, but also the kind of digital world we create for the next generation. My hope is that through collaborative yet flexible governance, we can harness the benefits of generative AI while safeguarding against its significant risks.